


# Apache Cassandra

The goal is to create Apache Cassandra clusters of about 5 nodes each. Each cluster will be run by one team and each team member will run one cluster node.



Please choose your team members now.

# Apache Cassandra

- Boot up your CIP-Pool Host and choose Ubuntu (not Windows!)
- Download Apache Cassandra "Latest GA Version" as tar.gz
- Start a terminal (e.g. **konsole**). Tip: Switch to a monospaced font.
- Unpack it into your home directory (tar xvfz filename.tar.gz)
- Determine the lowest and highest IP address in your team. This will be the **seed hosts**
- Edit the file conf/cassandra.yaml.
  - Comment out the **listen\_address** directive
  - Comment in the **listen\_interface** directive and set it to **eno1**
  - Set the value of the **seeds** key to the IPs of the seed hosts determined earlier
  - Agree on a cluster name and set the value of the **cluster\_name** key accordingly
- Change into the **bin/** directory and start the **./cassandra** script one node after another, starting with the seed nodes. Wait at least a minute between starts.
- Watch out for error messages

# Apache Cassandra

## Checking the cluster:

- Open another terminal window or tab
- Change into the **bin/** directory
- Start the script **./nodetool status**
- The output should look something like:

```
Datacenter: datacenter1
=====
Status=Up/Down
|/ State=Normal/Leaving/Joining/Moving
--  Address            Load           Tokens       Owns (effective)  Host ID
UN  172.20.38.87        79.2 KiB       16           60.3%             720a6a7e-2615-4f77-8e4a-652870ea
UN  172.20.38.88        140.66 KiB    16           60.9%             54b90b8d-ab37-4abe-b1eb-1da24ec7
UN  172.20.38.85        79.28 KiB     16           60.0%             6f0a4343-97ed-49c5-8255-b9f94cab
UN  172.20.38.86        79.2 KiB      16           58.7%             e12d7735-4b94-4207-8c8c-d6529306
UN  172.20.38.89        199.63 KiB    16           60.1%             5956d51f-9ab4-448e-aa2b-abf9c55e
```

- You might have to wait a bit. If it still doesn't work after some time:
  - Check if your java processes still run (ps ax), restart cassandra if not
  - Watch out for error messages
  - Cry for help

# Apache Cassandra

Change into the **bin/** directory and fire up the cli client with **./cqlsh**

Create a keyspace (similar to a database)

Every team member can create his own keyspace, use different names

Use a replication factor of 3 at first

The syntax is:

```
CREATE KEYSPACE <name>
  WITH REPLICATION = {
    'class' : 'SimpleStrategy',
    'replication_factor' : <factor>
  };
```

# Apache Cassandra

Now create a table and insert some data. The syntax is similar to SQL.

Example:

```
CREATE TABLE people (id INT PRIMARY KEY, name TEXT);  
  
INSERT INTO people (id, name) VALUES (1, 'Annie');  
INSERT INTO people (id, name) VALUES (2, 'Bernhard');  
INSERT INTO people (id, name) VALUES (3, 'Charlotte');  
INSERT INTO people (id, name) VALUES (4, 'David');  
INSERT INTO people (id, name) VALUES (5, 'Esmeralda');  
...
```

Insert at least 10 rows

Discuss with your team mates, how the data is probably distributed

# Apache Cassandra

Some handy commands:

cqlsh: CONSISTENCY (ONE|TWO|THREE|QUORUM|ALL) -- sets the consistency level for all following commands

Linux shell: killall -(STOP|CONT) java -- pause or continue the local node

Linux shell: ./nodetool stopdaemon -- stop the local node

Linux shell (in bin/): ./cassandra -- start the local node

Linux shell (in bin/): ./nodetool status -- check the cluster status

# Apache Cassandra

- Set the consistency level ALL
- Now stop the two cluster nodes with the highest IP addresses
- Ask for some of the rows by their id (works like an SQL select)
- Try to explain what you observe
- What changes with a consistency level of TWO? (first discuss, then try)

# Apache Cassandra

- Set the consistency level to ALL again, leave the two nodes stopped.
- Try to insert some data
- Use different primary keys for your attempts
- Again, try to explain what you observe



# Apache Cassandra

- We will try to simulate a network partition now:
- First, we have to find out some data records, that is also located on the two stopped nodes
- Set the consistency level to TWO, leave the two nodes stopped.
- Try to select random data records by id, until an error occurs
- Now set the consistency level to ONE and update the record by inserting a different name for the same id

# Apache Cassandra

- Stop the three running nodes now (wait until they have really shut down), then start the other two nodes again.
- Set the consistency level to TWO and try to retrieve the ID you have updated. What happens?
- Set the consistency level to ONE and try again. Which value do you get?
- Now start the three stopped nodes one by one, after each start try to see what happens to the value on the node of the previous step.
- Discuss what has happened

# Apache Cassandra

- There is a nodetool-function how to find out, which nodes contain a specific data record by primary key. Find out which one this is and test it. Make sure the information is correct by using CONSISTENCY and shutdown nodes.
- Develop an own test scenario to test replication or sharding behaviour of Apache Cassandra. Perform the test and evaluate the results.