

**Prof. Dr. Florian Heinz**  
[florian.heinz@sysv.de](mailto:florian.heinz@sysv.de)


# Modern Database Concepts

## The Map-Reduce Algorithm



OSTBAYERISCHE  
TECHNISCHE HOCHSCHULE  
REGENSBURG

# Map Reduce




- Processing Schema for very big amounts of data
- Gained popularity in 2004 through Google
- Paper: "MAPREDUCE: SIMPLIFIED DATA PROCESSING ON LARGE CLUSTERS" by Dean & Ghemawat, 2004

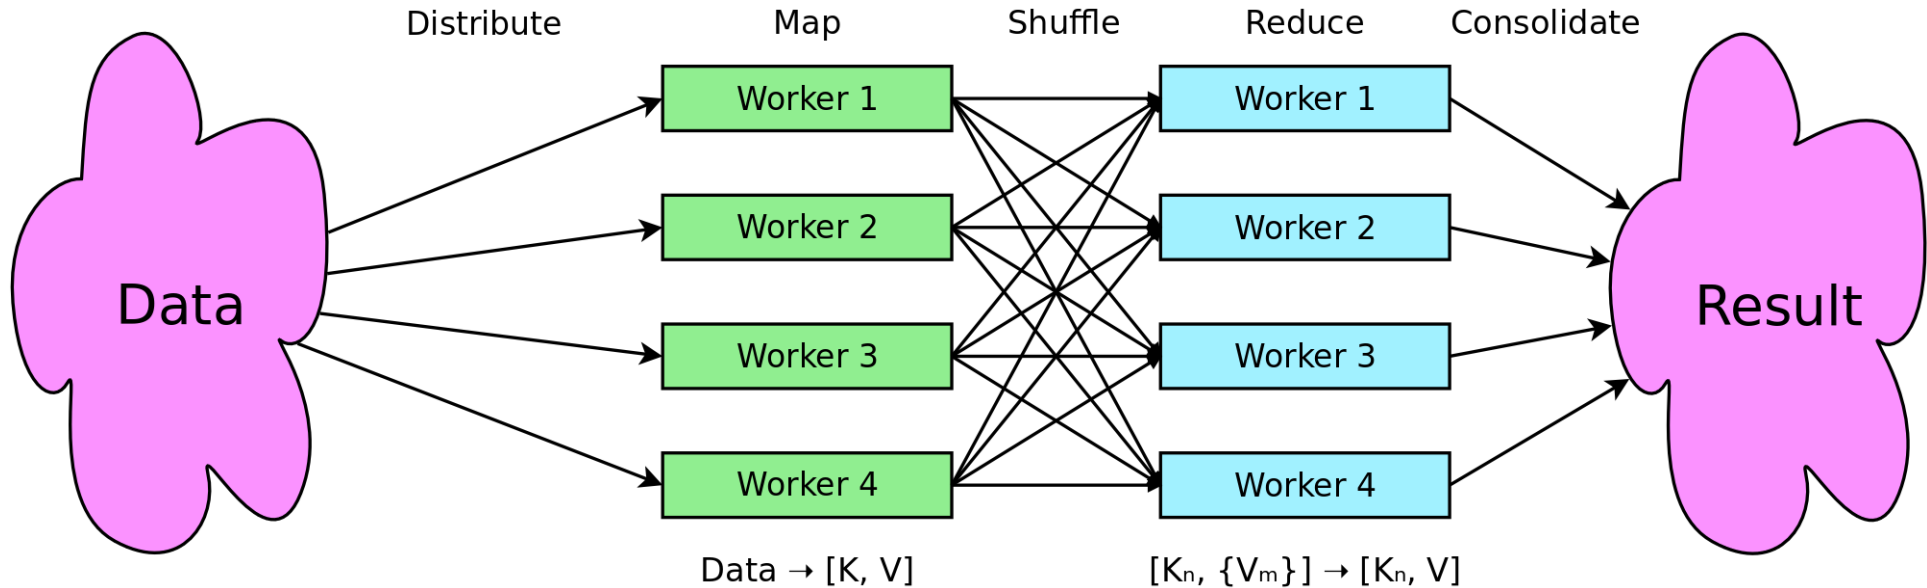
<https://static.googleusercontent.com/media/research.google.com/en//archive/osdi04.pdf>

- Useful for computations that can be massively parallelized


# Map Reduce




## Overview:



# Map Reduce: Data




- Data format can be arbitrary, no special structure required
- Example: Raw text files, XML documents, Images, Audio Files
- Suitable for data in **Data Lakes**
- Drawback: Map and Reduce functions have to be programmed at low level, e.g. in Java



# Running Example: Data

Input: CSV files with personal data (e.g. 100 files with many GB each)

```
# name,gender,birthdate,income
Anna Smith,female,1994-05-25,52000
Peter Miller,male,1996-07-02,45000
... <1 million lines>
```




Task: Calculate the average yearly income per age of the people


Similar to SQL Query:

```
SELECT age(birthdate) AS age, avg(income) AS income FROM persons
GROUP BY age
```

# Map Reduce: Distribute




- Distribution of the data items is usually handled by the framework
- A single large input file is split up into individual blocks on different nodes:
  - Sharding for load distribution ...
  - ... but also replication for fault tolerance
- A function to split the input into individual data items has to be provided by the user
- Examples: split based on words, lines or (binary) records
- The resulting items will be processed in parallel



# Running Example: Distribute

Input: CSV files with personal data (e.g. 100 files, 1 GB each)

```
# name,gender,birthdate,income  
Anna Smith,female,1994-05-25,52000  
Peter Miller,male,1996-07-02,45000  
... <100 files, each with 1 million records>
```




Output: Individual data items (e.g. 100 mio. items)

Large files are usually already split up in blocks on several nodes on the distributed storage


Distribution phase: Files are split up into individual records (CSV-lines), that go into the map phase.

MapReduce framework should try hard to process the data blocks locally.

# Map Reduce: Map




- First real processing step
- Map function has to be provided by the user
- Maps a data item into a Key/Value pair
- Massively parallel execution
  - Most of the work should happen in this step
- Data mapped to the same key is related and processed together in the reduce step





# Running Example: Map



Input: CSV lines with personal data (e.g. 100 million records)


```
Anna Smith,female,1994-05-25,52000 -> (27,52000)
Peter Miller,male,1996-07-02,45000 -> (25,45000)
... <100 million records>
```

Output: Key-Value pairs (e.g. 100 million pairs)


Each record is mapped to a Key/Value pair (age, income)

Parallel processing on many tasks and potentially distributed over the network

# Map Reduce: Shuffle




- Redistributes Key/Value pairs
- Handled by the MapReduce framework
- Preparation for reducing
- All pairs with the same key go to the same node
- Expensive operation: potentially high communication overhead
- optionally: Preceding **Combine** step to reduce node-local data prior to shuffling



# Running Example: Shuffle

Average income per



Input: Key-Value pairs (e.g. 100 mio. pairs)


```
(27, 52000), (27, 43000), (27, 78000), (27, 30000) -> Worker2  
(25, 45000), (25, 34000), (25, 23000), (25, 41000) -> Worker4  
... <many key/value pairs with about 50 different keys>
```

Same ages are distributed to the same workers for reducing (calculating the average)


In this example: 100 mio. pairs with 50 keys are distributed to 4 worker nodes for 50 reduction tasks

An optional **Combine** step could pre-aggregate the local (age,income) pairs grouped by age on each node. The result would be  $\leq 200$  pairs of (age, (count,avgincome)) ( $\leq 50$  on each node). The additional count is important for the final reduce step to weigh the intermediate results correctly

# Map Reduce: Reduce




- A single reduce step processes all data with the same key
- Result: usually a single Key/Value-pair (optionally more than one)
- Level of parallelism depends on number of different keys
- Bordercase: If only one key is used, a single reduce step will have to process all data (no parallelism)



# Running Example: Reduce

Input: One Key/Value-set pair per reduce task from kv-pairs with the same key


```
Worker 2.1: (27, {52000, 43000, 78000, 30000}) -> (27, 50750)
Worker 4.1: (25, {45000, 34000, 23000, 41000}) -> (25, 35750)
... <many reduce tasks on all workers>
```




One reduce task gets all income values for a single age.

The reduce task calculates the average income and emits a key/value pair with (age, avgincome)


# Map Reduce: Consolidate



- Collects results from the reduce step
- Usually written to a common data object
- e.g.: file on a distributed storage
- After that, the MapReduce algorithm is finished.



# Running Example: Consolidate



Input: Key/Value pairs from the reduce tasks


```
27, 50750  
25, 35750  
... <many output records>
```

Output: File with the result records

The key/value pairs from the reduce tasks are written to a common storage, i.e. a file on a distributed filesystem

After the MapReduce algorithm has terminated, this file can be fetched to collect the results

# Other Examples:



Word Count: Determine frequency of words


- Map: `<word> -> (<word>, 1)`
- Reduce: `(<word>, {1,...}) -> (<word>, n)`

Create Inverted Index:

- Map: `<word> -> (<word>, <docid>)`
- Reduce: `(<word>, {doc1,doc2,doc1,doc3,...}) -> (<word>, [doc1,doc2,doc3,...])`




# Map Reduce: Conclusion



## Advantages:

- Massive parallel computations
- Works on arbitrary data
- User only has to provide Map and Reduce functions



## Problems:

- Computation effort must significantly outweigh the communications overhead
- Low level programming needed

# Apache Hadoop



A Distributed Computing framework

Suitable for **Data Lakes**

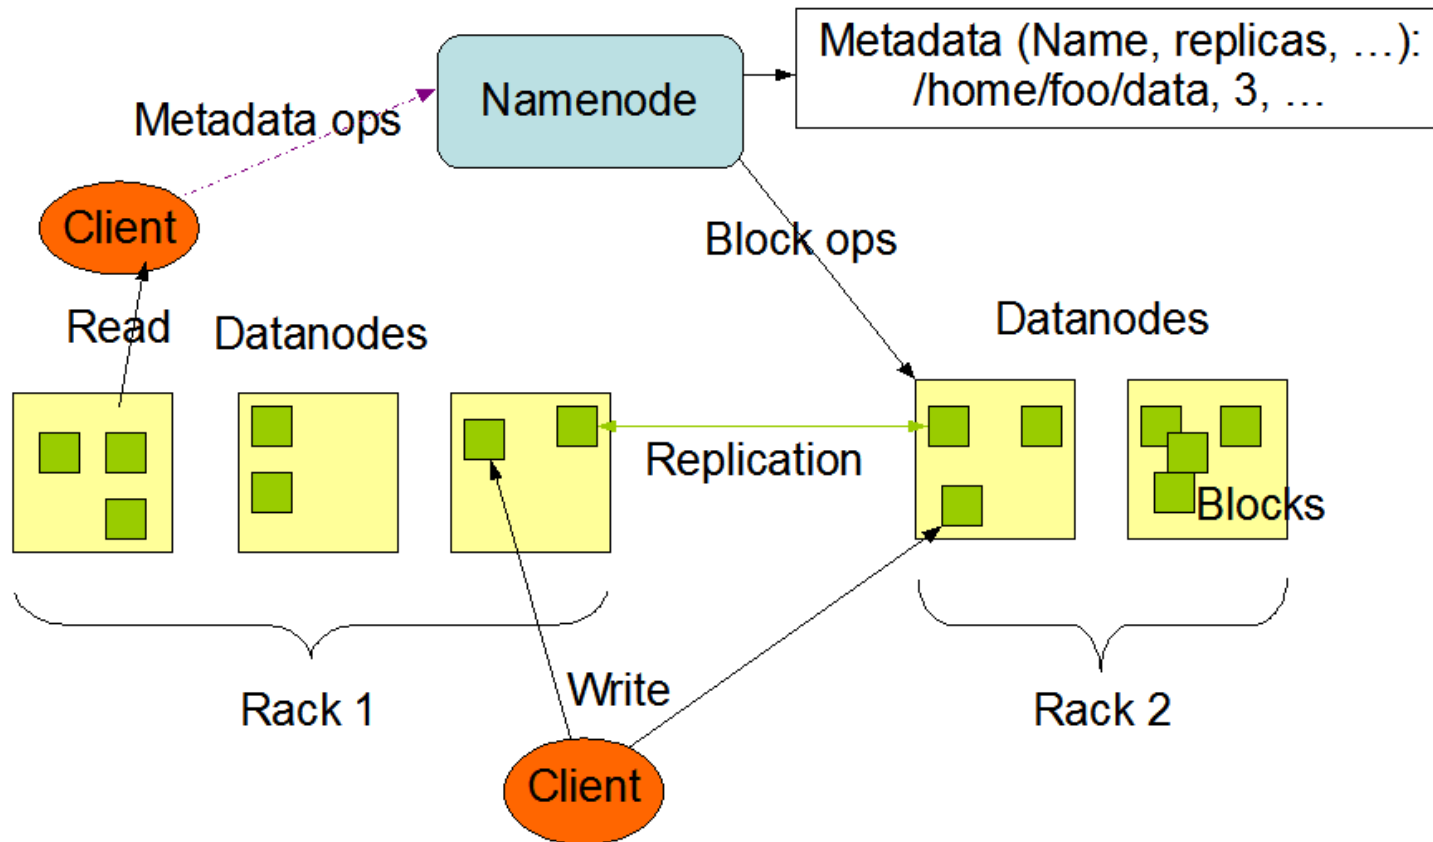
Base: HDFS (Hadoop Distributed File System)

```
$ bin/hdfs dfs -mkdir -p /ex1/input
$ bin/hdfs dfs -mkdir -p /ex1/output
$ bin/hdfs dfs -put *.csv /ex1/input
$ bin/hdfs dfs -ls /ex1/input
Found 6 items
-rw-r--r--  user  grp  38023  2022-05-01  21:34  /ex1/input/data1.csv
-rw-r--r--  user  grp  37998  2022-05-01  21:34  /ex1/input/data2.csv
-rw-r--r--  user  grp  37971  2022-05-01  21:34  /ex1/input/data3.csv
-rw-r--r--  user  grp  37913  2022-05-01  21:34  /ex1/input/data4.csv
-rw-r--r--  user  grp  37988  2022-05-01  21:34  /ex1/input/data5.csv
-rw-r--r--  user  grp  37993  2022-05-01  21:34  /ex1/input/data6.csv
```

# Apache Hadoop



## HDFS Architecture

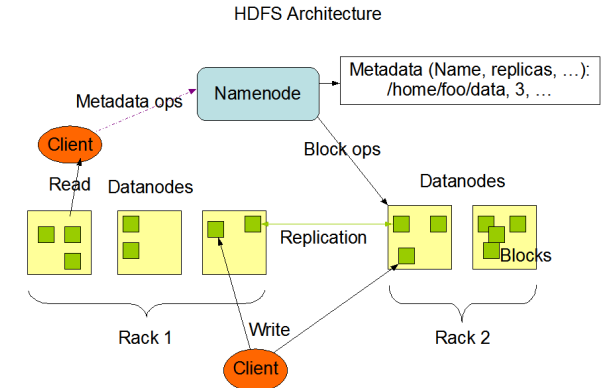


# Apache Hadoop



## Namenode:

- Stores Metadata of the File System
- Determines, which blocks go to which Datanode
- Client asks Namenode, where to find specific data.  
(e.g. bytes 15000-27000 of a specific file)
- Single point of failure, but high-availability with standby servers possible



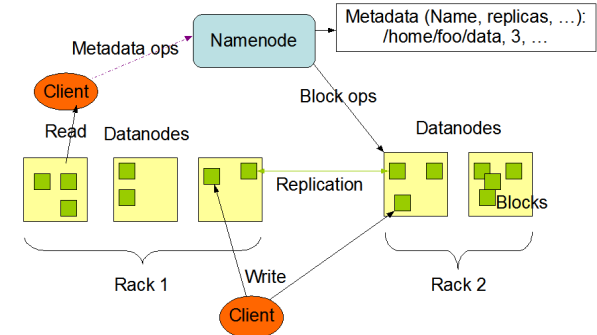
# Apache Hadoop



## Datanode:

- Stores the actual data
- Files are splitted into blocks of fixed size
- With replication factor  $n$ , each data block exists on  $n$  different nodes
- Client transmits data directly from/to datanodes

HDFS Architecture



More information: <https://hadoop.apache.org/docs/current/hadoop-project-dist/hadoop-hdfs/HdfsDesign.html>

# Apache Hadoop



MapReduce on input files

Generates output files

User has to provide java classes for:

- Input split function
- Map function
- Reduce function

# Apache Hadoop



Provides mechanisms for:

- Distributing data amongst the cluster
- Executing the map and reduce functions with minimal communication overhead
- Fault tolerance by replicating data shards and reassigning work to active nodes

# Apache Hadoop



## Example AvgIncome:

```
$ javac -cp $(bin/hadoop classpath) AvgIncome.java
$ jar -cf ai.jar AvgIncome*.class
$ bin/hadoop jar ai.jar AvgIncome /ex1/input /ex1/output
$ bin/hdfs dfs -cat /ex1/output/part-r-00000
27  48156
28  50481
[...]
42  50696
```



# Apache Hadoop



Map Function AvgIncome:

```
public static class CSVMapper
    extends Mapper<Object, Text, IntWritable, IntWritable> {

    public void map(Object key, Text value, Context context)
        throws IOException, InterruptedException {
    try {
        StringTokenizer itr = new StringTokenizer(value.toString());
        String name = itr.nextToken(",");
        String gender = itr.nextToken(",");
        LocalDate birthdate = LocalDate.parse(itr.nextToken(","));
        Integer income = Integer.parseInt(itr.nextToken(","));
        LocalDate now = LocalDate.now();
        Period period = Period.between(birthdate, now);

        IntWritable age = new IntWritable(period.getYears());
        IntWritable inc = new IntWritable(income);
        context.write(age, inc);
    } catch (Exception e) { }
    }
}
```

# Apache Hadoop



Reduce Function AvgIncome:

```
public static class AvgCalculateReducer extends
    Reducer<IntWritable,IntWritable,IntWritable,IntWritable> {

    public void reduce(IntWritable key,
        Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0, len = 0;
        for (IntWritable val : values) {
            sum += val.get();
            len++;
        }
        IntWritable result = new IntWritable(sum/len);
        context.write(key, result);
    }
}
```