

Secure Programming

Input Validation

Prof. Dr. Christoph Skornia
christoph.skornia@hs-regensburg.de

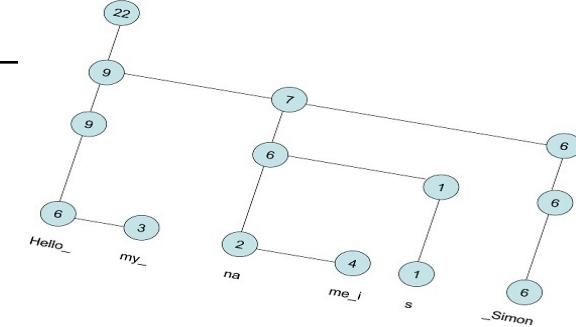
facts about strings:

- not built in type in C and C++
 - C: array of chars
 - C++: ISO/IEC98 is promoting `std::basic_string`, `std::string` and `std::wstring`
 - further structures like rope or cord for storing very long strings
- different definition for string-classes on different OS
 - `std::string` is UTF-8 by default
 - `std::wstring` is dependent on system and compiler
 - `std::wstring` is UTF-16 on Windows
 - `std::wstring` is UTF-32 on (typical) Unix-type systems (like Linux and OSX)
 - use `sizeof(wchar_t)` or list predefined macros (cpp -dM Ctrl+D) to clarify if unsure



facts about strings:

- null-termination
 - C-style strings are NULL-TERMINATED arrays
 - internal representation of `std::basic_string`, `std::string` and `std::wstring` does not need to be null terminated as well as Win32 `LSA_UNICODE_STRING`
- usually different types of strings are used (especially in legacy code)
 - different libraries require different types of strings
 - different type have different memory requirements
 - type conversions are necessary and may lead to vulnerabilities



Input Validation – Strings

common errors:

- Unbounded stream from stdin
bad example:

```
int main(void) {  
    char Password[80];  
    puts("Enter 8 character password:");  
    gets(Password);  
    ...  
}
```

better: limit size (truncate)

```
int main(void) {  
    char Password[9];  
    puts("Enter 8 character password:");  
    fgets(Password,9,stdin);  
    ...  
}
```

Input Validation – Strings

common errors:

- Unbounded string copy and concatenation example:

```
int main(int argc, char *argv[]) {
    char name[2048];
    strcpy(name, argv[1]);
    strcat(name, " = ");
    strcat(name, argv[2]);
    ...
}
```

better: dynamic allocation

```
int main(int argc, char *argv[]) {
    char *buff = (char *)malloc(strlen(argv[1])+1);
    if (buff != NULL) {
        strcpy(buff, argv[1]);
        printf("argv[1] = %s.\n", buff);
    }
    else {
        /* Couldn't get the memory – recover */
    }
    return 0;
}
```

Input Validation – Strings

common errors:

- Extracting characters from `cin` into a character array example:

```
int main(void) {
    char buf[12];

    cin >> buf;
    cout << "echo: " << buf << endl;
}
```

where is the problem?
what is `cin`?

`istream& operator>> (istream& is, char* str);`

- the problem is: no truncation!
- Solution: `(ios_base::width)`

```
#include <iostream>

int main(void) {
    char buf[12];

    cin.width(12);
    cin >> buf;
    cout << "echo: " << buf << endl;
}
```

Input Validation – Strings

common errors:

How many errors do you find?

```
1. int main(int argc, char* argv[]) {  
2.     char source[10];  
3.     strcpy(source, "0123456789");  
4.     char *dest = (char *)malloc(strlen(source))  
5.     int i;  
6.     for (i=1; i <= 11; i++) {  
7.         dest[i] = source[i];  
8.     }  
9.     dest[i] = '\0';  
10.    printf("dest = %s", dest);  
11.}
```

1. line 2 and 3: `source[10]` is 10 bytes long but `strcpy()` copies 11 byte into it
2. line 4: `malloc()` does not allocate memory for terminating null byte
3. line 5: index `i` starts with 1 although array indices start with 0
4. line 5: `i<=11` is counting one more than might be intended
5. line 8: terminating null will cause an out-of-bounds write

This type of errors is called “off-by-one” error.

Input Validation – Strings

Common errors:

- what is the output of:

```
int main(int argc, char* argv[]) {  
    char a[16];  
    char b[16];  
    char c[32];  
  
    strcpy(a, "0123456789abcdef");  
    strcpy(b, "0123456789abcdef");  
    strcpy(c, a);  
    strcat(c, b);  
    printf("a = %s\n", a);  
    return 0;  
}
```

answer: ask the compiler it
depends on the management
of stack memory



even worse: it depends on the content of
allocated memory prior to the start of the program.

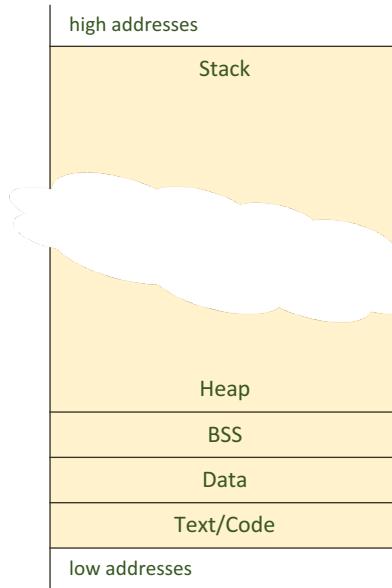
reason: null-termination mechanism, C-style-strings
and string-literals are null-terminated

Input Validation – Strings

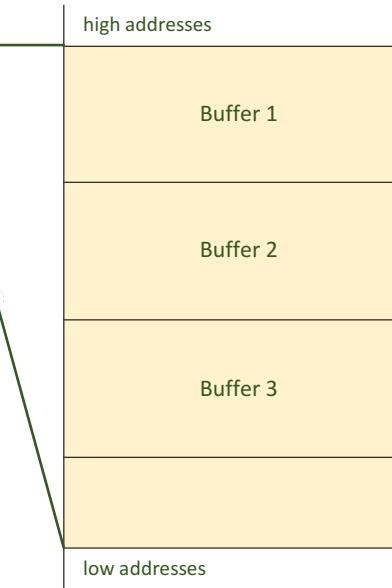
Vulnerabilities:

- ❑ What happens if you write more data to the stack than allowed?

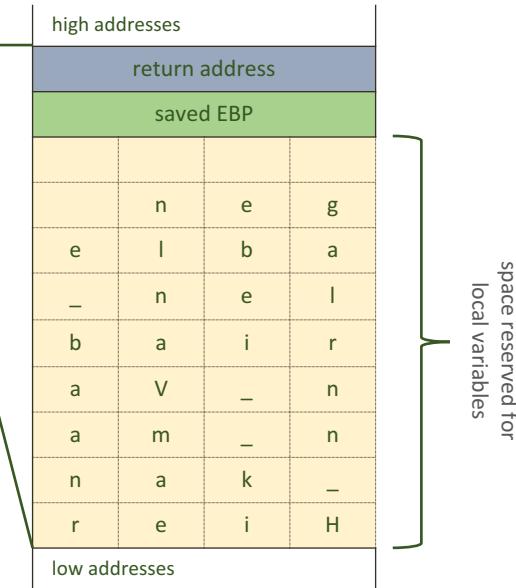
Memory Management



Stack :



Inside the Stack



Input Validation – Strings

- ❑ C and C++ simply don't care and continue to write to the memory
- ❑ consequence: the saved EBP and the return address will be overwritten

high addresses			
y	x	_	h
c	u	a	_
r	e	b	a
_	n	e	g
e	l	b	a
_	n	e	l
b	a	i	r
a	v	_	n
a	m	_	n
n	a	k	_
r	e	i	H
low addresses			

here was the return address

and here the saved EBP

- ❑ in case the return address is overwritten with an invalid value, the program will crash (this is already a successful denial of service (DOS) attack)
- ❑ in case the return address is overwritten with a valid value, the program will continue from there. This opens up a range of opportunities for targeted manipulation of the program.

Input Validation – Strings

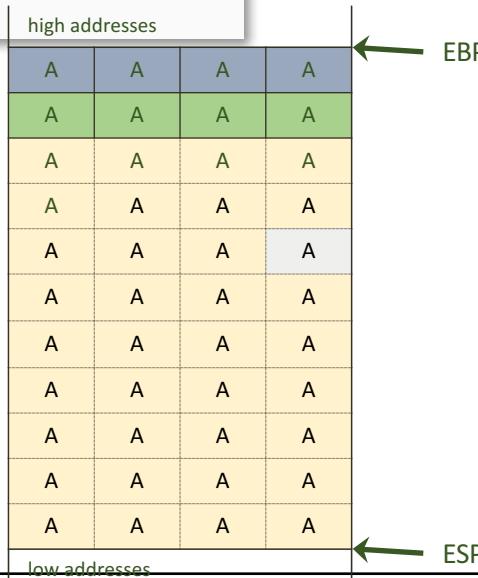
Smashing the stack (example 1)

```
#include <stdio.h>
#include <string.h>

int main(int argc, char** argv)
{
    char buffer [512];
    strcpy(buffer, argv[1]);
    printf("You entered:
%s\n",buffer);
    return 0;
}
```

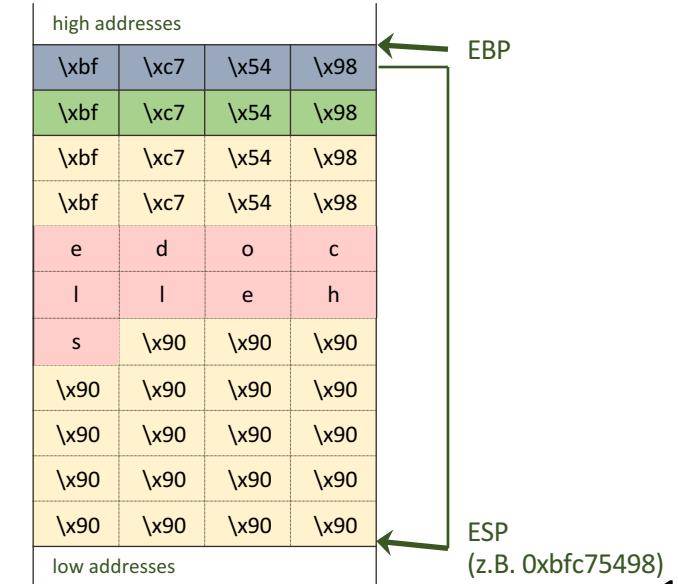
1. identify vulnerability

(by trying to crash the programm)



2. construct exploit

- a. NOPs (\x90)
(cpu simply moves forward to the next input)
- b. executable code (e.g. /bin/sh)
- c. return address which leads into the area of NOPs



□ Smashing the stack (example 2)

- What is this program doing?
- Is it vulnerable to DOS?
- Will the already shown stack-smashing attack work?
- Is there any option for a different type of attack?

```
05 void
06 geheim (void)
07 {
08     printf ("GEHEIM !!!\n");
09     exit (0);
10 }
11
12 void
13 oeffentlich (char *args)
14 {
15     char buff[12];
16
17     memset (buff, 'B', sizeof (buff));
18
19     strcpy (buff, args);
20     printf ("\nbuff: [%s] (%p)(%d)\n\n", &buff, buff, sizeof (buff));
21 }
22
23 int
24 main (int argc, char *argv[])
25 {
26     int uid;
27
28     uid = getuid ();
29
30     if (uid == 0)
31         geheim ();
32
33     if (argc > 1) {
34         printf ("geheim()    -> (%p)\n", geheim);
35         printf ("oeffentlich() -> (%p)\n", oeffentlich);
36
37         oeffentlich (argv[1]);
38
39         printf ("geheim() -> (%p)\n", geheim);
40         printf ("oeffentlich() -> (%p)\n", oeffentlich);
41     } else
42         printf ("Kein Argument!\n");
43
44 }
45 }
```

Input Validation – Strings

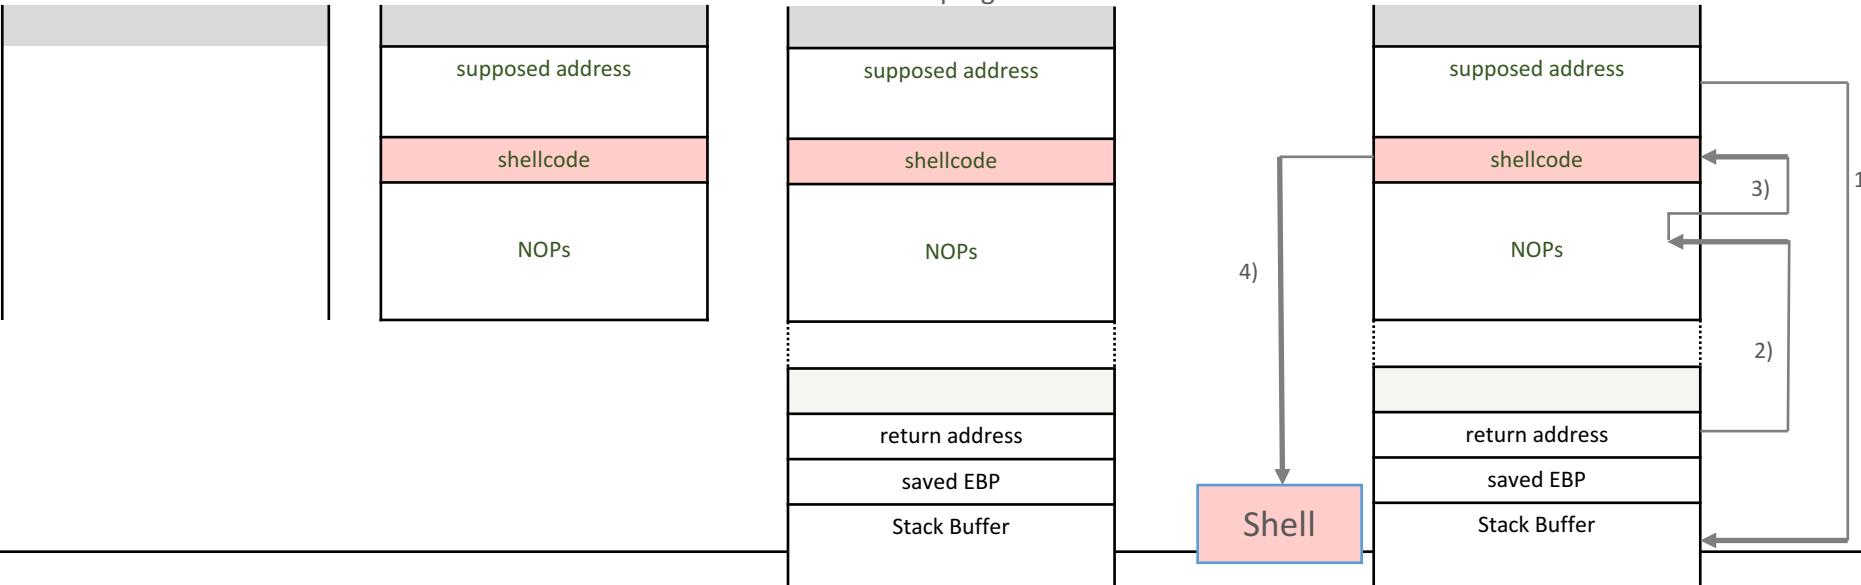
- ❑ In the last example 12 byte is (most likely) not enough space to place your code in
- ❑ What to do now?
- ❑ idea: save malicious code anywhere and write address to EIP
- ❑ how to do that?
- ❑ lets try environment variables.

prior to exploit

after exploit

after starting
vulnerable program

programm flow



Input Validation – Strings

□ Smashing the stack (example 3)

- What is this program doing?
- Is it vulnerable to DOS?
- Will the already shown stack-smashing attacks work?
- Is there any option for a different type of attack?
- Answer: Off by one error in line 10 allows potentially to overwrite last byte of EBP
- Most likely it does not work but would it be a problem?

```
01 #include <stdio.h>
02
03 static void
04 function (char *args)
05 {
06     char buff[256];
07     int i;
08
09
10    for (i = 0; i <= 256; i++)
11        buff[i] = args[i];
12 }
13
14 int
15 main (int argc, char *argv[])
16 {
17     if (argc < 2) {
18         printf ("No Argument!\n");
19         exit (1);
20     }
21
22     function (argv[1]);
23
24     return 0;
25 }
26 }
```

Input Validation – Strings

❑ Exploit: modification of last bit of saved ebp

Repetition: LEAVE and RET Function

LEAVE:

```
movl %ebp, %esp  
popl %ebp
```

RET:

```
call func  
movl $0, %eax eip now points here  
addl $8, %esp
```

high addresses			
\xbf	\xff	\xfa	\x34
\xbf	\xff	\xf9	\x80
A	A	A	A
e	d	o	c
I	I	e	h
s	\x90	\x90	\x90
\x90	\x90	\x90	\x90
\x90	\x90	\x90	\x90
\x90	\x90	\x90	\x90
low addresses			

⇒ after LEAVE of func:
esp: 0xbfffffa38 ebp:0xbfffffa34

⇒ after RET of func:
esp: 0xbfffffa40 ebp:0xbfffffa34

⇒ after LEAVE in main:
esp: 0xbfffffa38 ebp:0x41414141

⇒ after RET in main:
eip: 0xbffff980

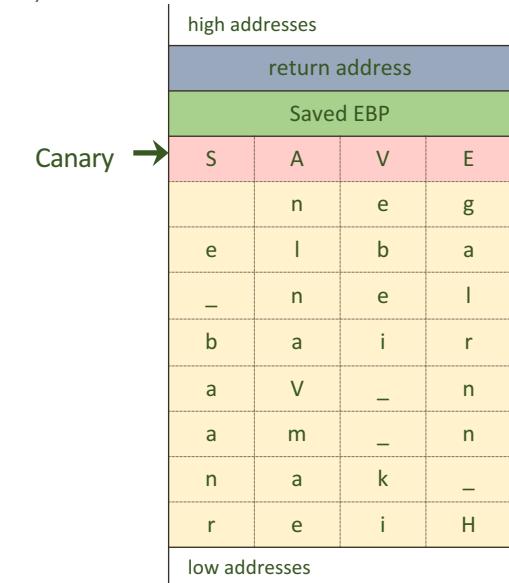
Modifying SFP will modify program flow

Attacks are possible e.g. OpenBSD FTP Daemon (2002)

Input Validation – Strings

□ Canaries

- include a "canary" value which, when destroyed, shows that a buffer preceding it in memory has been overflowed
 - Terminator canaries (NULL, CR, LF, -1)
problem: if canary is known, an attacker can try to work around it
 - Random canaries
generated at program initialization and stored in a global variable
(usually padded by unmapped pages, which makes it harder to read)
 - Random XOR canaries
Random Canaries that are XOR scrambled using control data
(change of control data a wrong canary)
- Implementations: StackGuard (Unix), StackCookie (Win)
- problem: DOS still working!

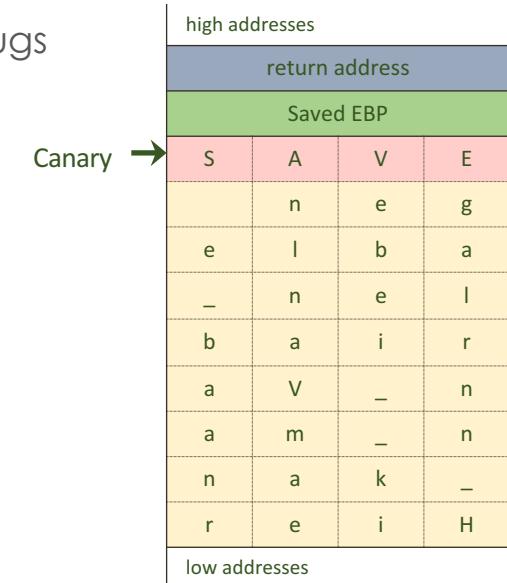


□ Bounds checking

- compiler-based technique that adds run-time bounds information for each allocated block of memory, and checks all pointers against those at run-time

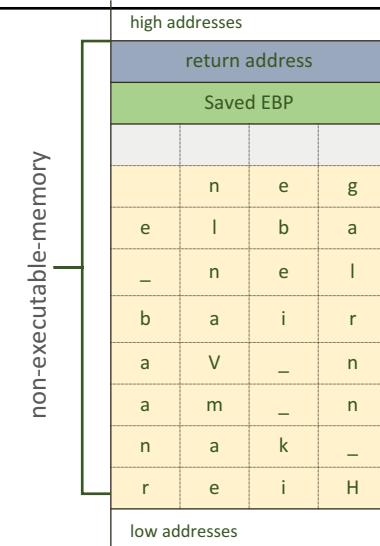
❑ AddressSanitizer (implemented in clang and gcc)

- ❑ compiler-based tool to detect all kinds of memory-corruption bugs
- ❑ `-fsanitize=address` oder `-fsanitize=kernel-address`
- ❑ performance overhead(average):
 - ❑ 73% cpu time
 - ❑ 340% memory usage
- ❑ Example features:
 - ❑ heap use after free
 - ❑ heap and Stack-buffer overflows



OS-based defense

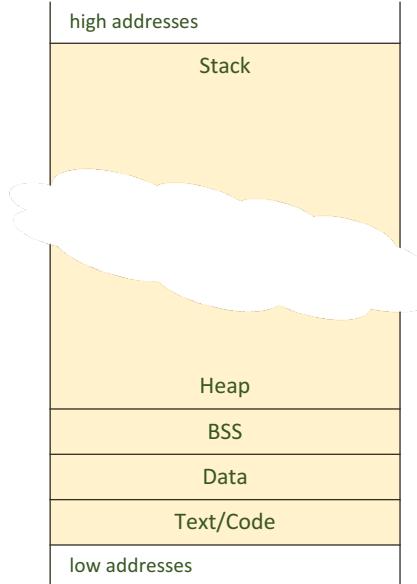
- ❑ Non-Executable Stack
 - ❑ mark data in stack as non-executable (-z no-exec-stack) in compiler
- ❑ Address-Space-Layer-Randomization (ASLR)
 - ❑ randomize address spaces
 - ❑ guessing of appropriate return addresses made difficult
 - ❑ Implementations: Standard in Unix and Win Systems
- ❑ Sandboxes and compartment based OS extensions
 - ❑ potentially vulnerable components (like services, special applications etc.) are executed in protected areas.
 - ❑ See Qubes OS, Ethos or new Kaspersky Press release



Input Validation – Strings

- What about this?

- Any vulnerabilities?



- How about memory management on the heap?

```
1 int main(int argc, char *argv[]){
2     FILE *filed;
3     char *userinput = (char *) malloc(20);
4     char *outputfile = (char *) malloc(20);
5
6     strcpy(outputfile, "/mynotes");
7     strcpy(userinput, argv[1]);
8
9     printf("userinput @ %p: %s\n", userinput, userinput);
10    printf("outputfile @ %p: %s\n", outputfile, outputfile);
11
12
13    filed = fopen(outputfile, "a");
14    if(filed == NULL)
15    {
16        fprintf(stderr, "error opening file %s\n", outputfile);
17        exit(1);
18    }
19
20
21    fprintf(filed, "%s\n", userinput);
22    fclose(filed);
23    return 0;
24 }
```

Input Validation – Strings

- What is this program doing?
- Any vulnerabilities?
- What would be the output of:
../fstring-2 ABCD_%1\\$x
- and what about:
../fstring-2 ABCD_%50\\$s

```
1 int main (int argc, char *argv[]){
2     int      a = 1;
3     int      b = 2;
4     int      c = 3;
5     char    buff1[32];
6     char    buff2[] = "BBB";
7     int      d = 4;
8
9     memset (buff1, 'A', 32);
10    snprintf (buff1, sizeof (buff1), argv[1]);
11    buff1[sizeof (buff1) - 1] = '\0';
12    printf ("buff1: [%s] (%d)\n", &buff1,strlen
13        (buff1));
14 }
```

Problem:

- `snprintf` is interpreting the format strings of the input

Consequence:

- read data from (normally not accessible) areas of the memory
- denial of service

Input Validation – Strings

- ❑ Any issues?

```
#include <stdio.h>
#include <string.h>

int
main (int argc, char *argv[])
{
    char user[512];
    strncpy(user, argv[1], 511);
    char outbuf[512];
    char buffer[512];
    sprintf(buffer, "ERR Wrong command: %.400s", user);
    sprintf(outbuf, buffer);
    return 0;
}
```

- ❑ What about

```
./fstring-5ABCD_%2000s
```

Buffer Overflow by buffer stretching!

Input Validation – Strings

- What does this code do?

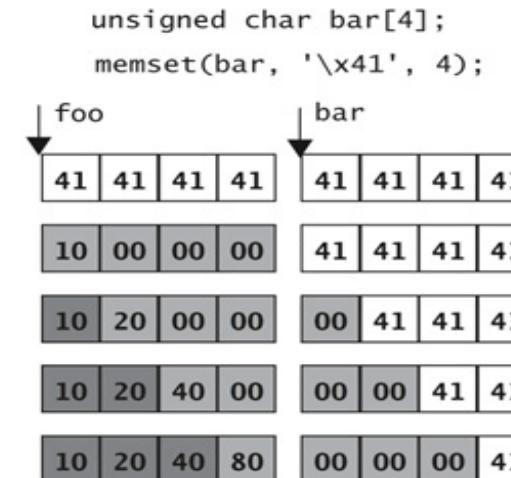
```
int i;  
printf("hello%n\n", (int *)&i);
```

- And this?

```
printf("\xdcl\xf5\x42\x01%08x%08x%08x%n");
```

Seems that it is possible to overwrite memory with that trick!

```
unsigned char foo[4];  
memset(foo, '\x41', 4);  
  
printf(  
    "%16u%n", 1, &foo[0]);  
printf(  
    "%32u%n", 1, &foo[1]);  
printf(  
    "%64u%n", 1, &foo[2]);  
printf(  
    "%128u%n", 1, &foo[3]);
```



Input Validation – Strings

Overwrite arbitrary address with
arbitrary value

- ❑ Lets look at the internal flow:

```
snprintf (buff1, sizeof (buff1), \x90\xf3\xff\xbf_%x_%x_%x_%n)
```

Format Strings:

- ❑ %x: Unsigned hexadecimal integer
- ❑ %n: Nothing printed.

The corresponding argument must be a pointer to a signed int.

The number of characters written so far is stored in the pointed location.

%n is overwriting the address at the first four bytes of buff1 (0xfffff390) with the length of buff1 (22)

- ❑ Even better

```
snprintf (buff1, sizeof (buff1), \x90\xf3\xff\xbf_%x_%x_%.10x_%n)
```

.10x is setting the length of the output and therefore the length of buff1

a Overwrite arbitrary address with arbitrary value!!!

Input Validation – Strings

Conclusion:

- ❑ Never use:
 - ❑ `printf (string);`
 - ❑ `fprintf (stderr, string);`
 - ❑ `snprintf (puffer, sizeof (puffer), string)`
- ❑ Better:
 - ❑ `printf ("%s", string);`
 - ❑ `fprintf (stderr, "%s", string);`
 - ❑ `snprintf (puffer, sizeof (puffer), "%s", string);`

Consider Compiler Warning options:

```
-fsyntax-only -pedantic -pedantic-errors -w -Wextra -Wall -Waddress -Waggregate-return -Warray-bounds -Wno-attributes -Wno-builtin-macro-redefined -Wc++-compat -Wc++0x-compat -Wcast-align -Wcast-qual -Wchar-subscripts -Wclobbered -Wcomment -Wconversion -Wcoverage-mismatch -Wno-deprecated -Wno-deprecated-declarations -Wdisabled-optimization -Wno-div-by-zero -Wempty-body -Wenum-compare -Wno-endif-labels -Werror -Werror=*= -Wfatal-errors -Wfloat-equal -Wformat -Wformat=2 -Wno-format-contains-nul -Wno-format-extra-args -Wformat-nonliteral -Wformat-security -Wformat-y2k -Wframe-larger-than=len -Wignored-qualifieds -Wimplicit -Wimplicit-function-declaration -Wimplicit-int -Winit-self -Winline -Wno-int-to-pointer-cast -Wno-invalid-offsetof -Winvalid-pch -Wlarger-than=len -Wunsafe-loop-optimizations -Wlogical-op -Wlong-long -Wmain -Wmissing-braces -Wmissing-field-initializers -Wmissing-format-attribute -Wmissing/include-dirs -Wmissing-noreturn -Wno-mudflap -Wno-multichar -Wnonnull -Wno-overflow -Woverlength-strings -Wpacked -Wpacked-bitfield-compat -Wpadded -Wparentheses -Wpedantic-ms-format -Wno-pedantic-ms-format -Wpointer-arith -Wno-pointer-to-int-cast -Wredundant-decls -Wreturn-type -Wsequence-point -Wshadow -Wsign-compare -Wsign-conversion -Wstack-protector -Wstrict-aliasing -Wstrict-aliasing=n -Wstrict-overflow -Wstrict-overflow=n -Wswitch -Wswitch-default -Wswitch-enum -Wsync-nand -Wsystem-headers -Wtrigraphs -Wtype-limits -Wundef -Wuninitialized -Wunknown-pragmas -Wno-pragmas -Wunreachable-code -Wunused -Wunused-function -Wunused-label -Wunused-parameter -Wunused-value -Wunused-variable -Wunused-but-set-parameter -Wunused-but-set-variable -Wvariadic-macros -Wvla -Wvolatile-register-var -Wwrite-strings
```

```
#include <stdio.h>
#include <string.h>

int main (int argc, char *argv[])
{
    int     a = 1;
    int     b = 2;
    int     c = 3;
    char   buff1[32];
    char   buff2[] = "BBB";
    int     d = 4;

    memset (buff1, 'A', 32);
    snprintf (buff1, sizeof (buff1), "%s", argv[1]);
    buff1[sizeof (buff1) - 1] = '\0';
    printf ("buff1: [%s] (%d)\n", &buff1, strlen (buff1));
}
```

Input Validation – Strings

Famous examples:

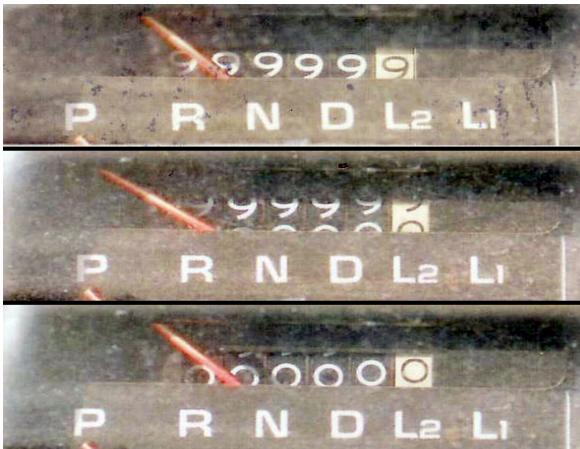
- [format string in Perl program](#)
- [bad call to syslog function](#)
- [bad call to syslog function](#)
- [NNTP server responses](#)
- [specifiers in a .bmp filename](#)
- [malicious internationalization messages](#)



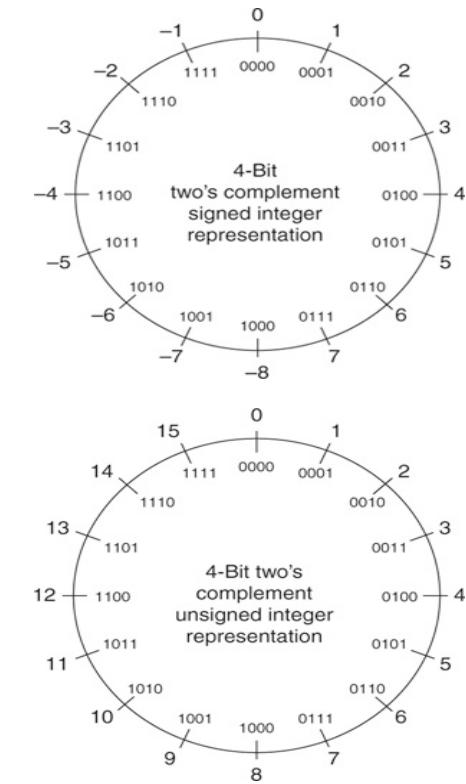
Input Validation – Integers

Regular errors with Integers:

- Integer Overflow ↗*e.g., Encryption uses very large integers*
- Sign Errors
- Truncation Error
try to know how large the integers can and will possibly get in your program



```
int main (){  
    int i;  
    unsigned int j;  
  
    i = 2147483647;  
    i++;  
    printf("i = %d\n", i);  
  
    j = 429967295;  
    j++;  
    printf("j = %u\n", j);  
  
    i = -2147483647;  
    i--;  
    printf("i = %d\n", i);  
  
    j = 0;  
    j--;  
    printf("j = %u\n", j);  
    i = -3;  
    unsigned short u;  
  
    u = i;  
    printf("u = %hu\n", u);  
  
    unsigned short int us = 32768;  
    short int is;  
  
    is = us;  
    printf("is = %d\n", is);  
  
    us = 65535;  
    is = us;  
    printf("is = %d\n", is);  
}
```



Input Validation – Integers

Corresponding Vulnerabilities:

- Integer Overflow
 - example (2000):
real world integer overflow in jpeg
comment fields
 - if comment field size is 1 then the
comment length calculates to 1-
 $2 = 0xffffffff$
 - memory allocation goes terribly wrong
 - example (2002):
computation of the memory region size by
calloc
 - <http://cert.uni-stuttgart.de/ticker/advisories/calloc.html>
- Similar Examples for sign and truncation
errors
 - <http://www.kb.cert.org/vuls/id/540517>
 - ... (ask google)

```
void getComment(unsigned int len, char *src) {  
    unsigned int size;  
  
    size = len - 2;  
    char *comment = (char *)malloc(size + 1);  
    memcpy(comment, src, size);  
    return;  
}  
  
int _tmain(int argc, _TCHAR* argv[]) {  
    getComment(1, "Comment ");  
    return 0;  
}
```



Input Validation – Integers

Corresponding Vulnerabilities:

- Any Vulnerability here?

What happens if pos=-4?

*negative position is
table & write into
memory addresses
you're not allowed to*

```
int *table = NULL;

int insert_in_table(int pos, int value){
    if (!table) {
        table = (int *)malloc(sizeof(int) * 100);
    }
    if (pos > 99) {
        return -1;
    }
    table[pos] = value;
    return 0;
}
```

pos value is passed as an argument, value would be written to a location
pos x 4 bytes before the start of the actual buffer!

Better:

```
int *table = NULL;

int insert_in_table(unsigned int pos, int value){
    if (!table) {
        table = (int *)malloc(sizeof(int) * 100);
    }
    if (pos > 99) {
        return -1;
    }
    table[pos] = value;
    return 0;
}
```

Input Validation – Integers

Mitigation Strategies:

- Range Checking
- Strong Typing
- Compiler Options (e.g. `-ftrapv` on gcc)
- Arbitrary Precision Arithmetic
 - GNU Multiple Precision Arithmetic Library (GMP)
 - Java BigInteger



- Validating Email-Addresses
 - check against RFC 822
 - default deny
- Cross Site Scripting
 - Refuse to accept anything that looks as if it may be HTML
 - Escape the special characters that enable a browser to interpret data as HTML
 - default deny
- SQL-Injection
 - Restrict user input to the smallest character set possible and refuse any input that contains characters outside of that set
 - Escape character that have special significance to SQL command processors
 - default deny

Default Deny!!!



thanks for your interest

to be continued

